

Vrije Universiteit Amsterdam



Bachelor Thesis

Co-Evolution of Generalist Morphology and Control for Diverse Environments

Author: Kubilay Tarhan (2721178)

1st supervisor: Anil Yaman
2nd reader: supervisor name

*A thesis submitted in fulfillment of the requirements for
the VU Bachelor of Science degree in Computer Science*

July 27, 2024

Abstract

This section should contain a brief summary of the key points of your paper, including the problem statement, methodology, results, and conclusion.

Keywords: keyword1, keyword2, keyword3

1 Introduction

The co-evolution of both morphology and control Since was first demonstrated by Karl Sims in his 1994 publication, "Evolving Virtual Creatures" [1], where he demonstrated virtual creatures interacting within a simulated three-dimensional physical environment. Since then, numerous researchers have pursued the simultaneous co-evolution of both morphology and control [2–5]. By co-evolving both morphology and control, researchers can develop more effective AI agent systems for various tasks. This approach leverages the principle of embodied cognition, which suggests that intelligence arises not solely from the brain or an agent’s control system, but from the dynamic interaction between the brain, body, and environment [6].

Moreover, co-evolving a morphology-controller pair (MC-pair) provides major benefits in robotic development, such as the ability to avoid the conventional physical design stage. Co-evolution enables to search for a potential morphological structures by evolving the MC-pair in a simulated environment, instead of wasting time iteratively re-designing and testing the robot in the physical world. Using an approach like this can furthermore speed up the development process and encourage the creation of more innovative morphologies that might not have been thought of in more traditional designs or initially deemed plausible.

Unfortunately, an evolved MC-pair will still specialize on the environment it has been trained in. This specialization occurs because the evolutionary process fine-tunes both the morphology and controller to the specific conditions of the training environment, resulting in high task performance in only the training environments but reduced task performance to new or marginally different environments. In the real-world, environments are inherently dynamic and unpredictable. Factors such as changing weather conditions, varying terrains, and unforeseen obstacles can significantly alter the operational context of an agent. This variability poses a substantial challenge to the robustness and gen-

eralizability of evolved MC-pairs. For example, an agent accustomed to smooth interior surfaces would find it difficult to maneuver over muddy, gravelly, or steeply inclined outdoor terrains. Agents created for static surroundings may also not be able to adjust to environments with changing lighting or moving impediments. Therefore, research into evolving a generalist MC-pair that can operate effectively across a wide range of environments is critically important.

To overcome this limitation, it is crucial to focus on developing robustness and generalizability in evolved agents. Robustness refers to the ability of an agent to maintain desirable behavior despite variations or perturbations in its input and output data [7–9]. Generalizability, on the other hand, refers to the ability of an agent to maintain desirable behavior under different conditions to those encountered during training [8, 9].

In this paper, we investigate the co-evolution of a generalist MC-pair for a wide range of environments, utilizing the Farama Gymnasium’s Ant-v4 task [10] as our testing framework. We modified it to allow for the evolution of both the leg lengths and widths of the ant, thereby adapting its morphology to suit various environments. Our training methodology employs an incremental variable training schedule, which is heavily based on the training algorithm specified by Triebold et al. [11], with minor modifications tailored to our objective of evolving a generalist MC-pair for a wide range of environments. **Our results show...**

2 Background Information

2.1 Co-optimizing of morphology and control

2.1.1 Embodied cognition

Despite the increase of computing power and knowledge, significant progress in the co-evolution of MC-pairs has remained stagnant. Researchers have proposed various hypotheses causing this, ranging from deficiencies in optimization algorithms to the notion that the training environments are not complex enough to facilitate morphological evolution, as outlined by Cheney et al. [12]. However, they suggest a new hypotheses to explain the difficulty in co-evolving MC-pairs, which is the theory of embodied cognition.

Embodied cognition suggests that the cognitive

processes arise not solely from the brain, but are a product of the dynamic interaction between the brain and morphology. This interplay means that even minor changes in morphology can disrupt the connection between the controller and morphology, requiring the controller to adapt itself again to the new morphology. This is also the reason why in [12] the morphologies generally converge before the 100th generation out of a total of 5000 generations, because at that point, further morphological evolution becomes less beneficial as it will only result in lower fitness scores. This phenomenon of premature convergence poses a challenge in co-evolving MC-pairs.

2.1.2 Premature convergence

Premature convergence occurs when the population quickly converges to a local optimum, resulting in a lack of genetic diversity and suboptimal solutions. In the co-evolution of MC-pairs, this leads to early stagnation of morphology evolution, as morphological changes disrupt optimized controllers and are being discarded by selection pressure, preventing the benefits of co-evolution [5].

For the co-evolution of MC-pairs, Lehman et al. [13] propose the Novelty Search with Local Competition (NSLC) algorithm, which utilizes a multi-objective search to optimize both diverse morphologies and fitness in conjunction with local competition, rewarding agents that outperform others with similar morphologies. While this approach did show an increase in the novelty of morphology, it did not yield higher fitness scores compared to using only a fitness objective function. Another method to address premature convergence is fitness sharing, which modifies the fitness function so that agents that are similar get penalized, thereby preserving and encouraging more morphological diversity [14]. Subsequent studies by Cheney et al. [2] propose explicitly protecting agents that have undergone a recent morphological mutation, which reduces the selection pressure and gives the controller more time to adapt to the new morphology. Using the open-source soft-body simulator VoxCad as the physics engine, their experiments showed that this method produced significantly higher fitness scores and increased morphological diversity. Additionally, it delayed premature convergence, as the best-performing agents emerged in later generations, who initially would have been discarded otherwise. Stensby et al. [3] extended upon this work by utilizing a more

indirect approach to mitigate premature convergence. They discuss that increasing the agents' exploration of new morphologies can be facilitated by training the agents in a diverse range of environments. These environments were made incrementally more challenging using the Paired Open-Ended Trailblazer (POET) algorithm. Their findings demonstrated that environments generated by POET increased morphology diversity, indicating that POET, or other forms of curriculum training, could be effective in delaying convergence.

2.2 Evolutionary strategies

2.2.1 Neuroevolution

Neuroevolution, the process of evolving artificial neural networks through evolutionary algorithms, has been proven to be a great alternative to traditional reinforcement learning approaches, particularly suited for tasks with continuous and high-dimensional input spaces [15]. This method begins by generating an initial population, and then selects the best solutions based on an evaluation function to produce the next generation of solutions. Evolutionary algorithms can be subdivided into genetic algorithms and evolutionary strategies. Genetic algorithms apply principles of evolution, such as selection, mutation, and recombination, typically encoding the parameters into binary strings. In contrast, evolutionary strategies utilize solely mutation and do not encode the parameters [16]. One of the challenges with these approaches is that they require an extensive number of evaluations due to the size of the population. Nonetheless, it is sometimes feasible to substitute the original evaluation function with a less accurate one, but lower in computational costs [17].

2.2.2 CPPN-NEAT

Conventionally, the topology of the neural network was established and fixed prior to training, which is a huge drawback, because the network's topology significantly influences the agent's performance. This drawback led to the development of Topology and Weight Evolving Artificial Neural Networks (TWEANNs), with the most prominent being the NeuroEvolution of Augmenting Topologies (NEAT) algorithm [18]. This algorithm evolves both the weights and the topology of neural networks. It initializes a population of minimal networks with only input and output nodes and incrementally adds nodes

and connections through mutations while utilizing a genetic algorithm for optimization. An important technique NEAT uses is speciation, which protects mutated networks, ensuring that new structures are not discarded and have the chance to evolve.

A notable advantage of NEAT is its use in conjunction with Compositional Pattern-Producing Networks (CPPNs). Unlike conventional neural networks, CPPNs utilize a variety of activation functions to generate complex and regular patterns [19]. This can be employed for encoding morphology parameters by mapping morphology features to specific values. In a study by Cheney et al. [20], CPPN-NEAT was used to encode the morphology of a soft robot, demonstrating that generative encoding using CPPNs assigned tissue cells to the voxels in a logical way, ensuring global coordination, resulting in improved locomotion. In contrast, the direct encoding assigned the voxels more independently of their neighboring voxels, resulting in less coordination and worsening locomotion.

2.2.3 Constraint-handling

Many optimization situations are subject to inherent constraints, such as the angular limits of a robotic arm, the allowed range of values for a mathematical function or simply in real life to use as less of a resource as possible. To effectively adhere to these constraints, evolutionary algorithms can employ several methods that direct the search away from prohibited areas of the solution space. For instance, penalty functions, which reduce the fitness score of solutions that violate constraints, therefore discouraging their selection. Additionally, repair algorithms actively modify solutions to meet the constraints that are in place. Another method is the multiobjective approach, which treats each constraint as a separate objective to be minimized alongside the main objective [21].

2.3 Robustness and generalizability

Robustness refers to the ability of an agent to maintain desirable behavior despite variations or perturbations in its input and output data [7–9]. Consequently, a robust agent is less susceptible to input and output perturbations. This makes robustness a critical attribute, as inputs and outputs from the training environment can differ significantly from those in the testing environment, especially in real-world scenarios. For instance, an input perturba-

tion can be caused by a sensor defect, resulting in slight measurement errors. In reinforcement learning, this means that we need to build robustness against the uncertainty of state observations and the actual state. Similarly, an output perturbation can result from a motory issue of the agent. In reinforcement learning, this means that we need to build robustness against uncertain actions between the actions generated by the agent and the conducted actions [9].

On the other hand, generalizability refers to the ability of an agent to maintain desirable behavior under different conditions to those encountered during training. It assumes correct input and output data and is concerned on the actual differences between the training environment and the testing or production environment. Testing generalization is divided into two distinctive parts. The first part is called interpolation, which requires the agent to perform well in environments similar to those of training, with both the testing and training environment parameters drawn from the same distribution. The second part is called extrapolation, which requires the agent to perform well in environments different from those of training, with the testing and training environment parameters drawn from separate distributions [8, 9].

There are some ways to achieve more robust agents. In this example [22], they compared the robustness of the models DENSER and NSGA-Net on the CIFAR-10 image classification task. It was found that the DENSER model exhibited an higher robustness, which concludes that certain architectures inherently have a higher degree of exhibiting robustness. Furthermore, one of the most popular method is adversarial training, where the agent is trained on adverserially perturbed training data. One way of doing this is finding the worst case perturbation at each training episode and training the model using the dataset with this perturbation [23].

The two main approaches to achieving generalizable agents are either training a generalist agent, which involves a trade-off in performance under specific conditions compared to a specialist agent, as shown by Triebold et al. [11], or developing agents that can explicitly adapt to certain conditions [8]. In this paper, we will focus on the first approach. Recently, there has been a growing recognition of the importance of using variability to train better and more generalized agents. Raviv et al. [24] discusses the relationship between vari-

ability and learning outcomes, highlighting a universal principle that variability enhances learning. This principle also holds true for machine learning, where employing a more variable learning schedule can improve an agent’s generalizability. Similarly, Stensby et al. [3], applied a similar principle of variability by training agents in a curriculum of environments generated by the Paired Open-Ended Trailblazer (POET) algorithm, where the generated environments were incrementally more difficult, enabling the agent to learn more efficient and complex behaviors. These studies demonstrate that using a variable learning schedule increases both robustness and generalizability.

3 Method

For the evolution of our MC-pair, we adopted the same algorithm initially proposed by Triebold et al. [11] with some minor modifications. Consider the set of training environments $E = \{e_1, e_2, \dots, e_n\}$ and the validation set $V = \{v_1, v_2, \dots, v_n\}$, where in our study $E = V$. Both E and V were established prior to training. Triebold et al. explored different training schedules, which determined in what order the training set was used during the evolutionary process. They found that using an incremental training schedule, where the environment is modified incrementally after every generation, yields the best results. Therefore, we will only consider the incremental training schedule in our analysis. Lastly, we initialize an empty generalist MC-pair set $G = \{\}$, which will store the evolved generalist MC-pairs throughout the evolutionary process, and an empty environment partition set $P = \{\}$, which will store partitions of the set V that correspond to the a generalist MC-pair in G . The algorithm will partition the environments and train a generalist MC-pair for that partition. This occurs when the environments differ significantly from one another and a single generalist MC-pair is not feasible.

The evolutionary process starts by initializing the first generation of MC-pairs, comprising both the ANN weights $\vec{W} = \{w_1, w_2, \dots, w_n\}$ and the morphology parameters $\vec{M} = \{m_1, m_2, \dots, m_n\}$. For the optimization process, \vec{W} and \vec{M} are concatenated into a singular vector $\vec{I} = \{w_1, w_2, \dots, w_n, m_1, m_2, \dots, m_n\}$ and fed to the XNES optimizer. Because the order of magnitude of the ANN and morphology parameters are

very different, we encode the morphology parameters to the same order of magnitude as the ANN parameters.

After every generation i , each MC-pair of the population is evaluated on the current training environment e_i , and the MC-pair with the highest fitness score, denoted as $I_{i, \text{best}}$, undergoes further evaluation on the entire validation set V , producing a generalist score g_{best} . If this generalist score is an improvement over MC_{best} , it is stored in the variable MC_{best} . After h number of generation, when no improvement is found, the evolutionary process is stagnated and the current MC_{best} is evaluated on the validation set, giving a $MC_{\text{best, mean}}$ and $MC_{\text{best, std}}$. Each environment in V , where MC_{best} scored higher than $MC_{\text{best, mean}} - MC_{\text{best, std}}$ will be added as a partition to P and removed from E and V . This process will repeat until E and V are empty, and thus every partition corresponds to a generalist MC-pair.

We integrated a penalty function within the fitness evaluation to ensure adherence to the constraints set for the morphological parameters. This function considers the decoded morphological parameters \vec{M} , alongside the lower and upper bounds of the constraints, C_{lb} and C_{ub} respectively, a scalar α , and a growth rate r^i , where i is the number of generations. The penalty function is defined as:

$$\text{Penalty} = \alpha r^i \cdot \sum (\max(0, C_{\text{lb}} - \vec{M}) + \max(0, \vec{M} - C_{\text{ub}})) \quad (1)$$

Thus the generalist fitness function then becomes:

$$g_{\text{best}} = \frac{1}{|V|} \sum_{i=1}^{|V|} (\text{evaluate}(MC_{\text{best}}, v_i) - \text{Penalty}) \quad (2)$$

Here $\text{evaluate}(MC_{\text{best}}, v_i)$ represents the evaluation function, returning the score of the best MC-pair of that generation on the validation environment v_i .

4 Experimental Setup

4.1 Environment

We use the Farama Gymnasium’s Ant-v4 environment [10] to evolve generalist MC-pairs. This environment consists of a flat plane for the ant to walk on, with the objective of traversing a long distance. The ant’s morphology comprises of four legs, each composed of an upper and a lower leg, with the upper leg attached to the torso. Figure 1 shows




Figure 1: Example of an ant in the environment. The red legs depict the lower leg part, and the light beige represents the upper leg part.

the appearance of the ant. To evolve a generalist MC-pair, we have to enable the ant to modify its morphology. Each leg part is parameterized by two attributes: length and width, resulting in a total of 16 distinct morphological parameters. The leg lengths are constrained between 0.1 and 1.5, while the widths are limited from 0.08 to 0.2.

To ensure a fast and successful experiment, additional modifications to the ant’s environment were necessary. Normally, the simulation would terminate when the ant’s torso ascends or descends to a specific height. This range has been increased to allow the ant to grow longer legs. Furthermore, additional settings were introduced to terminate the run whenever no significant forward movement was detected, suggesting that the ant had frozen in place. Another custom rule was implemented to detect if the ant flipped upside down by monitoring the z-vector of the torso. These additions significantly reduced evaluation times, especially at the start of the evolutionary process.

4.2 Experimental parameters

For the controller, a fully connected feedforward ANN is evolved, based on the topology used by Triebold et al. [11]. This topology consists of a single hidden layer with 20 neurons. The input layer corresponds to the ant’s continuous observation space, which includes 27 observations, such as the angles between the torso and leg connections

or velocity values, resulting in 27 input nodes. The output layer corresponds to the ant’s action space, where actions are values ranging from $[-1, 1]$, representing the torque applied to the rotors. This results in a total of 8 actions and thus 8 output nodes. Each layer also includes a bias node.

For the parameters used in the XNES algorithm, we used the default settings provided by the algorithm. The initial standard deviation of the search was set to 0.01, and the initial parameter ranges for the controller and encoded morphology parameters were $[-0.1, 0.1]$. To decode the morphology parameters, a linear transformation is used where the lower bound morphology constraint maps to -0.1 and the upper bound morphology constraint maps to 0.1. The lower and upper bound constraints used for the length and width of the legs are $[0.1, 1.5]$ $[0.05, 0.2]$ respectively. **Should I also insert these linear transformation formula’s that I used specifically?**

The parameters used for the penalty function were decided based on initial experimental analysis. For the growth rate r^i , we settled on 1.03 in conjunction with two scale factors for α . This means we are using two different penalty functions. The reason for this is that when the morphological parameters are decoded into negative values, the environment will crash because it cannot accommodate negative values for morphology. In these cases, α is set to 1000; otherwise, it is set to 100.

For the evolutionary process, we decided to evaluate the generalist scores after 2500 generations, instead of at every generation from the start. The reason for this is to decrease computational cost, but primarily because at the start of the evolutionary process, it is difficult to achieve high fitness scores due to the search for an effective morphology to exploit. After 2500 generations, a generalist score will be computed from the fittest individual in each generation. If there is no improvement found after 500 generations, the process is either terminated or a partition will occur.

Additionally, we conducted an alternative evolutionary run in which partitioning was disabled, allowing the run to extend to 5000 generations. The reason for this is because of the unpredictability of discovering a generalist MC-pair, due to it not being the objective function, but a secondary outcome of the method used. The exclusion of partitioning in this scenario will provide a valuable basis for comparison with the partitioned approach.

4.3 Training, evaluation and testing

The MC-pair’s training consists of different simulated environments, where the terrain is algorithmically generated based on the inputs. In this experiment, we consider three different environments. The first environment, referred to as the default environment, is a single static environment with a flat surface. The two remaining environments are dynamically generated, named the rough terrain environment and the hill terrain environment (see Figure 2 and Figure 3, respectively). The rough terrain environment is generated based on the block size and floor height. The block size determines the size of a block that can be elevated, and the floor height determines the range at which this block can be elevated. For the hill terrain environment, we utilize a Perlin noise map to determine the elevations of the terrain. The scale parameter determines the smoothness of the terrain, meaning that the terrain heights have smaller and subtler changes in height, leading to a more hilly terrain, with the floor height serving as the range at which these heights can be set.




Figure 2: Rough terrain environment generation, based on the block size and floor height.

In total, we have 40 variations of the rough terrain environment, 44 variations of the hill terrain environment, and the default environment, totaling up to a set of 85 different environments. We will split this set into training, validation, and testing sets. In our case, for the validation set, we will use the same set as for the training set. The training set contains the default environment, 24 hill environments, and 24 rough terrain environments, totaling to 49 environments. The testing set contains 20 hill environments and 16 rough terrain environments, totaling to 36 environments. The exact splits are shown in Table 1.

To test the generalizability of the MC-pair, we




Figure 3: Hills terrain environment generation, based on the scale of the Perlin noise and floor height.

run two separate tests. The first test is to evaluate the performance on environments similar to those of training, with both the testing and training environment parameters drawn from the same distribution, which we call interpolation. The second test is to evaluate the performance on environments different from those of training, with the testing and training environment parameters drawn from separate distributions, which we call extrapolation. In summary, to test interpolation, we will utilize the training set, and to test extrapolation, we will use the testing set described in Table 1. For each evolutionary run, all its data is derived by taking the median from five independent runs.

5 Results

Figure 4 shows the obtained fitness scores for each environment, represented in a heatmap. The by red enclosed cells indicates the environments from the testing set used for extrapolation testing, while the non-red-enclosed cells represent the environments from the training set used for interpolation testing.

For the hill environment, the heatmap shows high performance in the environments at the lower left and relatively lower performance when nearing the the upper right. This may indicate inherent complexity in certain environments compared to others. The fitness scores and their standard deviations on the training and testing sets are similar. The training set has a mean fitness score of 2997.99 with a standard deviation of 1175.13, and the testing set has a mean fitness score of 3012.55 with a standard deviation of 1183.09.

For rough terrain environment, the heatmap also shows high performance in the environments

Environment	Parameters	Training Set	Testing Set	Step
Rough terrain	Block size	[1, 4]	[1, 4]	1
	Floor height	[0.2, 0.4] \cup [0.7, 0.9]	[0.1] \cup [0.5, 0.6] \cup [1.0]	0.1
Hills terrain	Scale	[5, 20]	[5, 20]	5
	Floor height	[2.2, 2.6] \cup [3.4, 3.8]	[2.0] \cup [2.8, 3.2] \cup [4.0]	0.2

Table 1: This table shows the environment set splits and used ranges for the hill and rough environments.




Figure 4: Fitness heatmap from generalist MC-pair evolved over 5000 generations with partitions disabled

at the lower left and relatively lower performance when nearing the upper right, but also more at the lower right part. The fitness scores and their standard deviations on the training and testing sets are also similar. The training set has a mean fitness score of 2511.15 with a standard deviation of 1157.55, and the testing set has a mean fitness score of 2453.54 with a standard deviation of 1333.95.

6 Discussion

Interpretation of the result

Implications of the results

In our experimental design, the starting position of the ant was set to one unit above the maximum floor height of the environment. This adjustment ensured that the ant would not spawn into the ground. However, this configuration introduced a minor delay in the ant’s initial contact with the ground in environments with lower floor heights, thereby slightly disadvantaging them. However, this was not significantly noticable in the results, as the environments with lower floor heights are inherently easier than those with higher floor heights.

Discuss also the morphological evolution.

add visual graph for this. Another notable aspect of our experiment was the symmetry in the ant’s morphology, which rendered modifications to specific legs arbitrary. For instance, altering the length of legs 1 and 2 while shortening legs 3 and 4 results in the exact morphology to its reverse. This symmetry significantly reduced the number of novel morphologies the ant could evolve in drastically, despite a relatively large search space. To enhance the potential for more novel morphologies in future experiments, we could create and evolve the legs during the evolutionary process and include variations in the attachment points on the ant’s torso for the leg.

Given the diverse range of environments, each with unique challenges and complexities, employing TWEANNS might offer advantages over static ANN architectures. By using TWEANNS, it would be possible to evolve ANN structures specifically designed to generalize on all environments.

The methodology employed in this paper does not have an explicit objective function for generalizability; however, it emerges as a side effect of the applied methods. A future approach could involve making generalizability the primary objective. This approach would require increased com-

putational resources, because of the need to assign a generalist score across the entire population, as also highlighted by Triebold et al. [11].

Regarding the environment generation, each type of environment was generated with just two dimensions. Adding additional dimensions will increase the variability of these environments, potentially further improving generalizability. Beyond solely structural modifications to the environment, incorporating variables such as changes in gravity or variations in the drag force of the ground could provide more insights into the evolution of MC-pairs, better mimicking real-world scenarios. Additionally, investigating the step size variable, exploring its trade-offs and potentially identifying an optimal step size value, which could result in decreasing the size of the training sets, thereby decreasing the generalist score evaluation.

7 Conclusion

This work has presented a comprehensive investigation of the co-evolution of morphology and control, highlighting its benefits, such as eliminating the need to design the morphology prior to training and the potential to evolve more innovative and performant agents. However, it also addresses the challenges associated with evolving MC-pairs, such as premature convergence and embodied cognition. We examined the possibility of evolving an MC-pair that will generalize across a wide range of environments. To achieve this, we employed the XNES evolutionary strategy, where the morphology and ANN parameters are both optimized simultaneously, in conjunction with an incremental and variable training schedule.

References

- [1] Karl Sims. “Evolving virtual creatures”. In: *Proceedings of the 21st Annual Conference on Computer Graphics and Interactive Techniques*. SIGGRAPH ’94. New York, NY, USA: Association for Computing Machinery, 1994, pp. 15–22. ISBN: 0897916670. DOI: 10.1145/192161.192167. URL: <https://doi.org/10.1145/192161.192167>.
- [2] Nick Cheney et al. *Scalable Co-Optimization of Morphology and Control in Embodied Machines*. 2017. arXiv: 1706.06133 [cs.AI].
- [3] Emma Hjellbrekke Stensby, Kai Olav Ellefsen, and Kyrre Glette. “Co-optimising robot morphology and controller in a simulated open-ended environment”. In: *Applications of Evolutionary Computation* (2021), pp. 34–49. DOI: 10.1007/978-3-030-72699-7_3.
- [4] Joshua E. Auerbach and Josh C. Bongard. “Environmental Influence on the Evolution of Morphological Complexity in Machines”. In: *PLOS Computational Biology* 10 (2014), pp. 1–17. DOI: 10.1371/journal.pcbi.1003399. URL: <https://doi.org/10.1371/journal.pcbi.1003399>.
- [5] Luis Eguiarte-Morett and Wendy Aguilar. “Premature convergence in morphology and control co-evolution: a study”. In: *Adaptive Behavior* 32 (2024), pp. 137–165. DOI: 10.1177/10597123231198497. URL: <https://doi.org/10.1177/10597123231198497>.
- [6] Josh C. Bongard. “Evolutionary robotics”. In: *Commun. ACM* (2013), pp. 74–83. ISSN: 0001-0782. DOI: 10.1145/2493883. URL: <https://doi.org/10.1145/2493883>.
- [7] R. Mangal, A. V. Nori, and A. Orso. “Robustness of Neural Networks: A Probabilistic and Practical Approach”. In: *2019 IEEE/ACM 41st International Conference on Software Engineering: New Ideas and Emerging Results (ICSE-NIER)*. 2019, pp. 93–96. DOI: 10.1109/ICSE-NIER.2019.00032. URL: <https://doi.ieeecomputersociety.org/10.1109/ICSE-NIER.2019.00032>.
- [8] Charles Packer et al. *Assessing Generalization in Deep Reinforcement Learning*. 2019. arXiv: 1810.12282 [cs.LG].
- [9] Mengdi Xu et al. *Trustworthy Reinforcement Learning Against Intrinsic Vulnerabilities: Robustness, Safety, and Generalizability*. 2022. arXiv: 2209.08025 [cs.LG].
- [10] Mark Towers et al. *Gymnasium*. 2023. DOI: 10.5281/zenodo.8127025. URL: <https://gymnasium.farama.org/>.
- [11] Corinna Triebold and Anil Yaman. *Evolving generalist controllers to handle a wide range of morphological variations*. Sept. 2023. DOI: 10.13140/RG.2.2.27217.71522.

- [12] *On the Difficulty of Co-Optimizing Morphology and Control in Evolved Virtual Creatures*. Vol. ALIFE 2016, the Fifteenth International Conference on the Synthesis and Simulation of Living Systems. Artificial Life Conference Proceedings. 2016, pp. 226–233. DOI: 10.1162/978-0-262-33936-0-ch042. URL: <https://doi.org/10.1162/978-0-262-33936-0-ch042>.
- [13] Joel Lehman and Kenneth Stanley. “Evolving a diversity of creatures through novelty search and local competition”. In: 2011, pp. 211–218. DOI: 10.1145/2001576.2001606.
- [14] Robert I McKay. “Fitness Sharing in Genetic Programming.” In: *GECCO*. 2000, pp. 435–442.
- [15] Paolo Pagliuca, Nicola Milano, and Stefano Nolfi. “Efficacy of Modern Neuro-Evolutionary Strategies for Continuous Control Optimization”. In: *Frontiers in Robotics and AI* 7 (2020). ISSN: 2296-9144. DOI: 10.3389/frobt.2020.00098. URL: <https://www.frontiersin.org/journals/robotics-and-ai/articles/10.3389/frobt.2020.00098>.
- [16] Mehrdad Dianati, In-Soo Song, and Mark Treiber. “An Introduction to Genetic Algorithms and Evolution”. In: 2002. URL: <https://api.semanticscholar.org/CorpusID:10975919>.
- [17] Judith Echevarrieta, Etor Arza, and Aritz Pérez. *Speeding-up Evolutionary Algorithms to solve Black-Box Optimization Problems*. 2024. arXiv: 2309.13349 [cs.NE]. URL: <https://arxiv.org/abs/2309.13349>.
- [18] Kenneth O. Stanley and Risto Miikkulainen. “Evolving Neural Networks through Augmenting Topologies”. In: *Evolutionary Computation* 10.2 (2002), pp. 99–127. DOI: 10.1162/106365602320169811.
- [19] Kenneth O. Stanley. “Compositional pattern producing networks: A novel abstraction of development”. In: *Genetic Programming and Evolvable Machines* 8.2 (2007), pp. 131–162. ISSN: 1389-2576. DOI: 10.1007/s10710-007-9028-8. URL: <https://doi.org/10.1007/s10710-007-9028-8>.
- [20] Nick Cheney et al. “Unshackling evolution”. In: *ACM SIGEVOlution* 7 (2014), pp. 11–23. DOI: 10.1145/2661735.2661737.
- [21] Oliver Kramer. “A Review of Constraint-Handling Techniques for Evolution Strategies”. In: *Applied Comp. Int. Soft Computing* 2010 (Jan. 2010). DOI: 10.1155/2010/185063.
- [22] Inês Valentim, Nuno Lourenço, and Nuno Antunes. *Adversarial Robustness Assessment of NeuroEvolution Approaches*. 2022. arXiv: 2207.05451 [cs.NE].
- [23] Kai Liang Tan et al. *Robustifying Reinforcement Learning Agents via Action Space Adversarial Training*. 2020. arXiv: 2007.07176 [cs.LG].
- [24] Limor Raviv, Gary Lupyan, and Shawn Green. “How variability shapes learning and generalization”. In: *Trends in Cognitive Sciences* 26 (May 2022). DOI: 10.1016/j.tics.2022.03.007.